

RAWDATA

Section 1

SQL part 3

Programming & Advanced features

Henrik Bulskov & Troels Andreassen

Programming & Advanced features

- ❑ Accessing SQL From a Programming Language
 - JDBC, ODBC and ADO.NET
 - ADO.NET with C# will be covered in more detail in section two
- ❑ Programming the database
 - Functions and Procedural Constructs in SQL
 - Triggers in SQL
- ❑ Recursive Queries
- ❑ Advanced aggregation (Ranking, Windowing)
- ❑ Data Analysis and OLAP

JDBC and ODBC and ADO.NET

- ❑ API (application-program interface) for a program to interact with a database server
- ❑ Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- ❑ JDBC (Java Database Connectivity)
 - works with Java
- ❑ ODBC (Open Database Connectivity)
 - works with C, C++, C#, and Visual Basic
- ❑ ADO.NET
 - works with the .NET framework
 - will be used with C# on RAWDATA
 - especially with what's called Entity-Framework and LINQ
 -

JDBC

- ❑ **JDBC** is a Java API for communicating with database systems supporting SQL.
- ❑ JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- ❑ Model for communicating with the database:
 - 1) Open a connection
 - 2) Create a “statement” object
 - 3) Execute queries using the Statement object to send queries and fetch results
 - 4) Extract data from result set
 - 5) Close connection
 - (Use exception mechanism to handle errors)

```
//STEP 1. Import required packages
```

```
import java.sql.*;
```

Java & JDBC Code

```
public class FirstExample {
```

```
    static final String DB_URL = "jdbc:postgresql://localhost:5432/university"; // a JDBC url
```

```
    //static final String DB_URL
```

```
    static final String USER = "postgres";
```

```
    static final String PASS = "toor";
```

```
    public static void main(String[] args) {
```

```
        Connection conn = null;
```

```
        Statement stmt = null;
```

```
        try{
```

```
            //STEP 1: Register JDBC driver
```

```
            Class.forName("org.postgresql.Driver");
```

```
            //STEP 2: Open a connection
```

```
            System.out.println("Connecting to database...");
```

```
            conn = DriverManager.getConnection(DB_URL,USER,PASS);
```

```
            //STEP 3: Execute a query
```

```
            System.out.println("Creating statement...");
```

```
            stmt = conn.createStatement();
```

```
            String sql;
```

```
            sql = "SELECT id, name, salary FROM instructor";
```

```
            ResultSet rs = stmt.executeQuery(sql);
```

```
            //STEP 4: Extract data from result set
```

```
            while(rs.next()){
```

```
                //Retrieve by column name and display values
```

```
                System.out.println("ID: " + rs.getString("id") + " " + rs.getString("name") + " " + rs.getInt("salary"));
```

```
            }
```

```
Connecting to database...
```

```
Creating statement...
```

```
ID: 10101 Srinivasan 65000
```

```
ID: 12121 Wu 90000
```

```
ID: 15151 Mozart 40000
```

```
ID: 22222 Einstein 95000
```

```
ID: 32343 El Said 60000
```

```
ID: 33456 Gold 87000
```

```
ID: 45565 Katz 75000
```

```
ID: 58583 Califieri 62000
```

```
ID: 76543 Singh 80000
```

```
ID: 76766 Crick 72000
```

```
ID: 83821 Brandt 92000
```

```
ID: 98345 Kim 80000
```

```
Goodbye!
```

```
//STEP 5: Clean-up environment
```

```
    rs.close();
```

```
    stmt.close();
```

```
    conn.close();
```

```
}catch(Exception e){
```

```
    //Handle errors
```

```
    e.printStackTrace();
```

```
}
```

```
    System.out.println("Goodbye!");
```

```
//end main
```

```
//end FirstExample
```

ADO.NET

- ❑ The ADO.NET API provides functions to access data similar to the JDBC functions.
- ❑ Thus ADO.NET allows access to results of SQL queries
- ❑ A similar model for communicating with the database:
 - 1) Open a connection
 - 2) Create a “statement” object
 - 3) Execute queries using the Statement object to send queries and fetch results
 - 4) Extract data from result set
 - 5) Close connection

C# & ADO.NET

```
using System;
using Npgsql;

namespace AdoExample
{
    class Program
    {
        static void Main(string[] args)
        {
            var connString = "Host=localhost;Username=troels;Password=troels;Database=uni";

            using (var conn = new NpgsqlConnection(connString))
            {
                conn.Open();
                // Retrieve all instructors
                using (var cmd = new NpgsqlCommand("SELECT id, name, salary FROM instructor", conn))
                using (var rdr = cmd.ExecuteReader())
                {
                    while (rdr.Read()){
                        Console.WriteLine("ID: {0} {1} {2} \n", rdr[0], rdr[1], rdr[2]);
                    }
                }
            }
        }
    }
}
```

```
ID: 10101 Srinivasan 65000
ID: 12121 Wu 90000
ID: 15151 Mozart 40000
ID: 22222 Einstein 95000
ID: 32343 El Said 60000
ID: 33456 Gold 87000
ID: 45565 Katz 75000
ID: 58583 Califieri 62000
ID: 76543 Singh 80000
ID: 76766 Crick 72000
ID: 83821 Brandt 92000
ID: 98345 Kim 80000
```

Press any key to continue...

Functions and Procedural Constructs in SQL

Procedural Extensions and Stored Procedures

- ❑ SQL provides a **module** language
 - Permits definition of functions and procedures in SQL
- ❑ Functions
 - write your own functions and add them to the database
 - use them like any function predefined by the DBMS, that is, within expressions
- ❑ Stored Procedures
 - store procedures in the database
 - execute them by "calling" them from applications or interfaces to the DBMS
 - this permits external applications to operate on the database without knowing about internal details
 - you can, for instance, make your own dedicated API that provides functionality but hides the database structure
- ❑ Triggers
 - you can add special procedures that are executed automatically by the system as a side effect of a modification to the database

Procedural Extensions and Stored Procedures

❑ PostgreSQL specialities

- PostgreSQL provides probably the most advanced framework and language extension for adding functions and procedures to the DBMS
- PostgreSQL does not include an explicit notion of Stored Procedure, but (as in other languages) you can consider a function that does not return anything to be a Procedure
 - Stored Procedure ~ Function of type void
- most DBMS provide a **call** statement to execute a stored procedure,
- PostgreSQL allow Stored Procedures (void functions) and functions to be invoked by SELECT-expressions or by using a special **perform** command)

Functions and Procedures

- ❑ Since SQL:1999 the standard supports functions and procedures
 - Functions/procedures can be written in SQL itself, or in an external programming language.
 - Some database systems (including PostgreSQL) support a particularly useful construct:
 - **table-valued function**, (returning a relation as a result).
- ❑ SQL:1999 also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment, and others
- ❑ Many databases have proprietary procedural extensions to SQL that differ from SQL:1999.

SQL Functions

- ❑ Define a function.

```
create function hello (s char(20))  
returns char(50)  
begin  
return concat('hello, ',s,'!');  
end;
```

- ❑ Use the function.

```
select hello('world') 'Message to all';
```

```
create function hello (s char(20))  
returns char(50) as  
$$  
begin  
return concat('hello, ',s,'!');  
end;  
$$  
language plpgsql;
```

\$\$ is used to enclose
the body as a literal

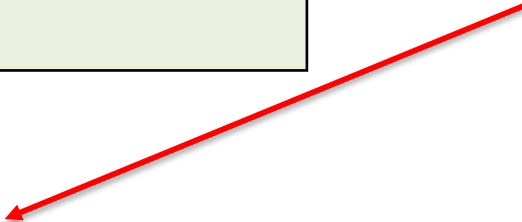
not using \$\$ would make “;” ambiguous

```
uni=# select hello('world') as "Message to all";  
Message to all  
-----  
hello, world!  
(1 row)
```

SQL Functions

```
create function hello (s char(20))  
returns char(50) as  
$$  
begin  
return concat('hello, ',s,'!');  
end;  
$$  
language plpgsql;
```

Same function but
now used on a table



```
uni=# select hello(name) "Message to all" from instructor;  
      Message to all  
-----  
hello, Srinivasan!  
hello, Wu!  
hello, Mozart!  
hello, Einstein!  
hello, El Said!  
hello, Gold!  
hello, Katz!  
hello, Califieri!  
hello, Singh!
```

SQL Functions

- ❑ Define a function that, given the name of a department, returns the **count of the number of instructors in that department**.

```
create function dept_count (dept_name varchar(20))  
returns integer  
begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```

DSC Figure 5.5

- ❑ Find the department name and budget of all departments with more than 1 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name) > 1
```

SQL Functions

- ❑ Same function, but now using the PL/pgSQL language
- ❑ Again **count of the number of instructors in that department.**

```
create function dept_count (d_name char(20))
returns integer as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$
language plpgsql;
```

```
uni=# select dept_count( 'Physics' );
dept_count
-----
          2
(1 row)
```

- ❑ or use the dept_count()-function in a where-condition.

```
create function dept_count (d_name char(20))
returns integer as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$
language plpgsql;
```

```
uni=# select dept_name, budget from department
uni-# where dept_count(dept_name ) > 1;
 dept_name | budget
-----+-----
Comp. Sci. | 100000.00
Finance    | 120000.00
History    | 50000.00
Physics    | 70000.00
(4 rows)
```


- ❑ or use the dept_count()-function in the select clause

```
create function dept_count (d_name char(20))
returns integer as $$
declare d_count integer;
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
    return d_count;
end;
$$
language plpgsql;
```

```
uni=# select distinct dept_name, dept_count(dept_name)
uni-# from department;
```

dept_name	dept_count
Biology	1
Comp. Sci.	3
Elec. Eng.	1
Finance	2
History	2
Music	1

SQL Procedures

❑ The *dept_count* function could instead be written as procedure:

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_name;
```

```
end
```

DSC page 175

❑ The SQL standard suggests that

- Procedures can be invoked either from an SQL procedure or
- from embedded SQL, using the **call** statement.

❑ Postgres

- does not include a Procedure construct
- does not include a call statement, but do have a perform
- only functions with return types can be defined
- however return type void would correspond to a procedure

SQL Procedures

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_name;
```

```
end
```

DSC page 175

- ❑ Since this procedure includes a value in the out-parameter *d_count* Postgres requires the definition to be an integer- rather than an void-function

```
create function dept_count_proc (in d_name varchar(20),  
                                out d_count int)  
  
returns integer as  
$$  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
end;  
$$  
language plpgsql;
```

SQL Procedures

calling the procedure from the command line

```
uni=# select dept_count_proc ('Physics');
 dept_count_proc
-----
                2
(1 row)
```

```
create function dept_count_proc (in d_name varchar(20),
                                out d_count int)
returns integer as
$$
begin
    select count(*) into d_count
    from instructor
    where instructor.dept_name = d_name;
end;
$$
language plpgsql;
```

- ❑ Nothing is gained from defining the out-parameter here
A more straight forward version would be

```
create function dept_count_proc (d_name varchar(20))  
returns integer as  
$$  
begin  
    return (select count(*)  
            from instructor  
            where instructor.dept_name = d_name);  
end;  
$$  
language plpgsql;
```

```
create function dept_count_proc (in d_name varchar(20),  
                                out d_count int)  
returns integer as  
$$  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name = d_name;  
end;  
$$  
language plpgsql;
```

SQL Procedures

```
create function dept_count_proc (d_name varchar(20))
returns integer as
$$
begin
    return (select count(*)
            from instructor
            where instructor.dept_name = d_name);
end;
$$
language plpgsql;
```

❑ or an even simpler as a plain SQL language function

```
create function dept_count_proc (d_name varchar(20))
returns bigint as
$$
    select count(*)
    from instructor
    where instructor.dept_name = d_name;
$$
language sql;
```

Table Functions

- ❑ SQL:2003 added functions that return a relation as a result

- ❑ Example: Return all instructors of a given department

create function *instructors_of* (*dept_name* **char**(20)

returns table (*ID* **varchar**(5),
 name **varchar**(20),
 dept_name **varchar**(20),
 salary **numeric**(8,2))

DSC Figure 5.6

return table

(**select** *ID, name, dept_name, salary*
from *instructor*
where *instructor.dept_name = instructors_of.dept_name*)

- ❑ Usage

select *
from table (*instructors_of* ('Physics'))

Table Functions

- ❑ Same function, but now using the SQL language (in the body)
- ❑ Again **used to retrieve instructors in the Physics department.**

```
create function instructors_of (dept_name char(20))
    returns table (ID varchar(5),
                  name varchar(20),
                  dept_name varchar(20),
                  salary numeric(8,2)) as
$$
    (select ID, name, dept_name, salary
     from instructor
     where instructor.dept_name = instructors_of.dept_name);
$$
language sql;
```

```
uni=# select * from instructors_of ('Physics');
   id   |   name   | dept_name |  salary
-----+-----+-----+-----
 22222 | Einstein | Physics  | 95000.00
 33456 | Gold     | Physics  | 87000.00
(2 rows)
```


Procedural Constructs

- ❑ Conditional statements (**if-then-else**)
- ❑ Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statement
- ❑ Loops (among others): **While** and **repeat** statements :
declare n integer default 0;
while $n < 10$ do
 set $n = n + 1$;
end while;

repeat
 set $n = n - 1$;
until $n = 0$
end repeat;
- ❑ Warning: most database systems implement their own variant of a modular (procedural) language – only inspired by the standard syntax

SQL Procedure, example with WHILE loop

```
drop table if exists foo;  
create table foo  
(  
    id serial primary key,  
    val integer  
);
```

a table, foo, for testing

defining the procedure

```
create or replace function load_foo()  
returns void as  
$$  
declare  
    i_max integer := 4;  
    i integer := 0;  
    n integer;  
begin  
    while i < i_max loop  
        n:=(random() * 10000);  
        insert into foo (val) values (n);  
        i:=i+1;  
    end loop;  
end  
$$  
language plpgsql;
```

calling and showing the effect

```
uni=# select load_foo();  
load_foo  
-----  
(1 row)  
  
uni=# select * from foo;  
id | val  
---+-----  
1  | 3822  
2  | 5438  
3  | 8353  
4  | 9690  
(4 rows)
```

SQL Procedure, example (cont.)

❑ Notice SQL-details

- drop ... if exists ... (very useful in a script you want to run repeatedly)
 - `drop table if exists foo;`
- auto incrementing primary key
 - `id serial primary key,`
- declaration and initialization of variable
 - `i_max integer := 4;`
- while loop to do several DML-statements
 - `while i < i_max loop`
- random() between 0 and 1 to generate number between 0 and 9999
 - `n:=(random() * 10000);`

Calling a procedure from another

```
drop table if exists foo;  
create table foo  
(  
    id serial primary key,  
    val integer  
);
```

a table, foo, for testing

defining the procedure that calls the procedure

```
create or replace function test()  
returns void as  
$$  
begin  
    perform load_foo();  
    perform load_foo();  
end  
$$  
language plpgsql;
```

calling and showing the effect

```
uni=# select test();  
test  
-----
```

(1 row)

```
uni=# select * from foo;  
 id | val  
----+-----  
  1 | 8537  
  2 | 3004  
  3 | 6846  
  4 | 3274  
  5 | 6931  
  6 | 4464  
  7 | 8140  
  8 | 5080  
(8 rows)
```

SQL Procedure, simplified using a FOR loop

```
create or replace function load_foo()  
returns void as  
$$  
declare  
    i_max integer := 4;  
    i integer := 0;  
    n integer;  
begin  
    while i < i_max loop  
        n:=(random() * 10000);  
        insert into foo (val) values (n);  
        i:=i+1;  
    end loop;  
end  
$$  
language plpgsql;
```

```
create or replace function load_foo()  
returns void as  
$$  
begin  
    for i in 1..4 loop  
        insert into foo (val) values (random() * 10000);  
        i:=i+1;  
    end loop;  
end  
$$  
language plpgsql;
```

Cursor

❑ **cursor**

- is a control structure that enables traversal of rows in a table
- a cursor is declared by a query and the table to be traversed is the result of this query

❑ **declare**

- Before a cursor can be used it must be declared (defined).
- `declare cur1 cursor for select name,salary from instructor;`

❑ **open** – perform the query

- The cursor must be opened for use. This process actually retrieves the data using the previously defined SELECT statement.
- `open cur1;`

❑ **fetch** – get the next row from the table

- Individual rows can be fetched (retrieved) as needed.
- `fetch cur1 into a, b;`

❑ **close** – close the cursor (clean up)

- When done, the cursor must be closed.
- `close cur1;`

SQL Procedure using cursor, example

a table, **vip**, for testing

```
drop table if exists vip;
create table vip as
  select name, salary
  from instructor;
truncate vip;
```

```
create or replace function curdemo()
returns void as
$$
DECLARE
  rec record;
  cur1 cursor for select name,salary from instructor;
begin
  open cur1;
  loop
    fetch cur1 into rec;
    exit when not found;
    if rec.salary > 81000 then
      insert into vip
        values (rec.name,rec.salary);
    end if;
  end loop;
  close cur1;
end;
$$
language plpgsql;
```

calling and showing the effect

```
uni=# select curdemo();
curdemo
-----

(1 row)

uni=# select * from vip;
   name  | salary
-----+-----
 Wu      | 90000.00
 Einstein | 95000.00
 Gold     | 87000.00
 Brandt   | 92000.00
(4 rows)
```

SQL Procedure using cursor, example(cont.)

- ❑ Notice SQL-details
 - the 4 “using cursors”-issues to remember:
 - **declare, open, fetch, close**
 - conditional statement (fairly standard)
 - `if ... then ... end if;`
 - a very useful data type **record**:
 - `rec record;`
 - `...`
 - `fetch curl into rec;`
 - another loop construction
 - `loop ... exit when ... end loop;`

Yet another loop ...

- ❑ A very useful loop in PostgreSQL is the following

```
[ <<label>> ]  
FOR target IN query LOOP  
    statements  
END LOOP [ label ];
```

- ❑ DO-block
 - anonymous function
 - very useful for adhoc tasks and for testing expressions
- ❑ raise notice ...
 - very useful for testing expressions

testing with a DO block

```
uni=# do $$  
uni$# declare  
uni$#     rec record;  
uni$# begin  
uni$#     for rec in select name  
uni$#         from instructor  
uni$#     loop  
uni$#         raise notice '%', rec.name;  
uni$#     end loop;  
uni$# end;  
uni$# $$;  
NOTICE:  Srinivasan  
NOTICE:  Wu  
NOTICE:  Mozart  
NOTICE:  Einstein  
NOTICE:  El Said  
NOTICE:  Gold  
NOTICE:  Katz  
NOTICE:  Califieri  
NOTICE:  Singh
```

SQL using cursor (now implicit), example

a table, **vip**, for testing

-- vip(name, salary)

```
create or replace function curdemo2()  
returns void as  
$$  
declare  
    rec record;  
begin  
    for rec in select name,salary from instructor  
    loop  
        if rec.salary > 81000 then  
            insert into vip  
                values (rec.name,rec.salary);  
        end if;  
    end loop;  
end;  
$$  
language plpgsql;
```

calling and showing the effect

```
uni=# truncate vip;  
TRUNCATE TABLE  
uni=# select curdemo2();  
curdemo2  
-----  
  
(1 row)  
  
uni=# select * from vip;  
      name      | salary  
-----+-----  
Wu              | 90000.00  
Einstein        | 95000.00  
Gold            | 87000.00  
Brandt          | 92000.00  
(4 rows)
```

SQL using cursor (now implicit), ex. (cont.)

- ❑ Compare to the cursor example above
 - the loop is changed to
 - **for rec in** `select name,salary from instructor`
 - **loop**
 - ...
 - **end loop**
 - the cursor is replaced by the expression given as argument in the for loop
 - `for rec in` **`select name,salary from instructor`**
 - conceptually this is an implicit cursor

External Language Functions/Procedures

- ❑ SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- ❑ Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out count integer)
```

```
language C
```

```
external name ' /usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))
```

```
returns integer
```

```
language C
```

```
external name ' /usr/avi/bin/dept_count'
```

External Language Functions/Procedures

- ❑ Notice the PostgreSQL **CREATE FUNCTION** statement:

```
CREATE FUNCTION function_name(...)
RETURNS type AS
BEGIN
-- logic
END;
LANGUAGE language_name;
```

- ❑ By default, PostgreSQL supports **three languages**:
 - SQL, PL/pgSQL, and C.
- ❑ You can also load other procedural languages
 - e.g., Perl, Python, and TCL

External Language Routines (Cont.)

- ❑ Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power.

- ❑ Drawbacks
 - Code to implement function may need to be loaded into database system and executed in the database system's address space.
 - risk of accidental corruption of database structures
 - security risk, allowing users access to unauthorized data

Why use Stored functions and procedures?

- ❑ Stored functions and procedures (routines) can be particularly useful
 - When **multiple client applications** are written in different languages or work **on different platforms**, but **need to perform the same database operations**.
 - When security is paramount. **Banks**, for example, **use stored procedures** and functions for all common operations
 - In addition, you can store **libraries of functions and procedures** in the database server
 - Provide **improved performance**. Less information needs to be sent between the server and the client.
 - Tradeoff: **increase the load on the database server**

Triggers

Triggers

- ❑ A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.

- ❑ To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.

Trigger Example – Referential constraint

- ❑ E.g. *time_slot_id* is not a primary key of *timeslot*, so we cannot create a foreign key constraint from *section* to *timeslot*.
- ❑ Alternative: use triggers on *section* and *timeslot* to enforce integrity constraints
- ❑ Figure 5.8 in DSC book:

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from timeslot)) /* time_slot_id not present in timeslot */
begin
    rollback
end;
```

Will not work in MySQL
Rollback is not allowed

PostgreSQL Triggers

- ❑ To create a new trigger in PostgreSQL:
 - Create a trigger function using CREATE FUNCTION statement.
 - Bind this trigger function to a table using CREATE TRIGGER statement.
- ❑ Create the trigger function

```
CREATE FUNCTION trigger_function() RETURN trigger AS
```

- a function similar to an ordinary function,
- does not take any arguments
- has return **return type trigger**
- important variables: **OLD** and **NEW** represent the states of row in the table before or after the triggering event.

PostgreSQL Triggers

- ❑ Create the trigger
 - use the CREATE TRIGGER statement:

```
CREATE TRIGGER trigger_name {BEFORE | AFTER | INSTEAD OF} {event [OR ...]}  
ON table_name  
[FOR [EACH] {ROW | STATEMENT}]  
EXECUTE PROCEDURE trigger_function
```

- The event could be INSERT, UPDATE, DELETE or TRUNCATE.
- (BEFORE) or after (AFTER) event specifies the order of the trigger and the update
- INSTEAD OF is used only for views
- two kinds of triggers: row level trigger and statement level trigger,

Trigger Example – Referential constraint

❑ Figure 5.8 in DSC book does NOT work

- Rollback is not allowed in a trigger

❑ The following is an alternative

- The result is the same: an update with a time_slot_id not present in the time_slot table will not be allowed (and will thus be ignored)

❑ Figure 5.8 in DSC book:

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from time_slot)) /* time_slot_id not present in time_slot */
begin
    rollback
end;
```

Will not work in MySQL
Rollback is not allowed

```
create function timecheck() -- the trigger function
```

```
returns trigger as $$
```

```
begin
```

```
    if (new.time_slot_id not in (select time_slot_id from time_slot)) then
        raise exception 'time_slot_id is unknown';
```

```
    end if;
```

```
end; $$
```

```
language plpgsql;
```

```
create trigger timecheck_trig -- the trigger (calling the trigger function
```

```
before insert on section
```

```
for each row execute procedure timecheck();
```

```
uni=# insert into section values
```

```
uni-# ('BIO-301', '2', 'Winter', '2009', 'Painter', '514', 'I');
```

```
ERROR: time_slot_id is unknown
```

```
CONTEXT: PL/pgSQL function timecheck() line 4 at RAISE
```

- ❑ Notice SQL and PostgreSQL details
 - The example is a **before** rather than an **after** trigger
 - **if** (inside the block) replaces **when** (outside)
 - “**referencing new row as *nrow***” won’t work, but you can reference the new value simply with **new**
 - **new** can be used in **insert** and **update**-triggers
 - **old** can be used similarly in **delete** and **update**-triggers
 - **raise exception** will prevent the insert and return an error message

```
create function timecheck() -- the trigger function
returns trigger as $$
begin
    if (new.time_slot_id not in (select time_slot_id from time_slot)) then
        raise exception 'time_slot_id is unknown';
    end if;
end; $$
language plpgsql;

create trigger timecheck_trig -- the trigger (calling the trigger function
before insert on section
for each row execute procedure timecheck();
```

```
uni=# insert into section values
uni-# ('BIO-301', '2', 'Winter', '2009', 'Painter', '514', 'I');
ERROR:  time_slot_id is unknown
CONTEXT:  PL/pgSQL function timecheck() line 4 at RAISE
```

Trigger Example – Ad hoc constraint

- ❑ Company policy (insert on instructor trigger)
 - No new employments in high budget departments (≥ 90000)
 - New employees (instructors) must never have a salary greater than everybody else

```
drop trigger if exists instructorcheck_trig on section;
drop function if exists instructorcheck;
create function instructorcheck() -- the trigger function
returns trigger as $$
begin
    if (new.dept_name not in (select dept_name from department where budget < 90000)) then
        raise exception 'No no no, no new employees in the % department', new.dept_name;
    end if;
    if (new.salary > (select max(salary) from instructor)) then
        raise exception 'No no no, salary too high';
    end if;
end; $$
language plpgsql;

create trigger instructorcheck_trig -- the trigger
before insert on instructor for each row execute procedure instructorcheck();
```

```
uni=# insert into instructor values (12345, 'Wong', 'Finance', 80000);
ERROR:  No no no, no new employees in the Finance department
```

```
uni=# insert into instructor values (23456, 'Wang', 'History', 100000);
ERROR:  No no no, salary too high
```

Triggering Events and Actions in SQL

- ❑ Triggering event can be **insert**, **delete** or **update**
- ❑ Triggers can be activated before an event, which can serve as extra constraints. E.g. convert blank grades to null.

```
create function setnull() -- the trigger function
returns trigger as $$
begin
    if (new.grade = ' ') then
        new.grade := null;
    end if;
    RETURN NEW;
end;$$
language plpgsql;

create trigger setnull_trig -- the trigger
before update on takes for each row execute procedure
setnull();
```


T

```
create function setnull() -- the trigger function
returns trigger as $$
begin
    if (new.grade = ' ') then
        new.grade := null;
    end if;
    RETURN NEW;
end;$$
language plpgsql;

create trigger setnull_trig -- the trigger
before update on takes for each row execute procedure
setnull();
```

testing ...

```
uni=# select * from takes where grade is null;
```

id	course_id	sec_id	semester	year	grade
98988	BIO-301	1	Summer	2010	

(1 row)

```
uni=# update takes
```

```
uni=# set grade=' ' where id ='98765' and course_id='CS-101';
UPDATE 1
```

```
uni=# select * from takes where grade is null;
```

id	course_id	sec_id	semester	year	grade
98988	BIO-301	1	Summer	2010	
98765	CS-101	1	Fall	2009	

(2 rows)

Trigger to Maintain credits_earned value

❑ Figure 5.9 from the DSC book

– create trigger *credits_earned* after update of *takes* on (*grade*)
referencing new row as *nrow*
referencing old row as *orow*
for each row
when *nrow.grade* \neq 'F' and *nrow.grade* is not null
and (*orow.grade* = 'F' or *orow.grade* is null)
begin atomic
update *student*
set *tot_cred* = *tot_cred* +
 (select *credits*
 from *course*
 where *course.course_id* = *nrow.course_id*)
where *student.id* = *nrow.id*;
end;

Will not work in PostgreSQL

Trigger to Maintain credits_earned value (Cont.)

❑ Figure 5.9 from the DSC book (New PostgreSQL version)

```
create function credits_earned() -- the trigger function
returns trigger as $$
begin
    if (new.grade <> 'F' and new.grade is not null
        and (old.grade = 'F' or old.grade is null)) then
        update student
        set tot_cred = tot_cred +
            (select credits from course
             where course.course_id= new.course_id)
        where student.id = new.id;
    end if;
end;$$
language plpgsql;

create trigger credits_earned_trig -- the trigger
    after update on takes for each row execute procedure
credits_earned();
```

- **NOTICE:** “update of *takes* on (*grade*)” is not supported in Postgres
- But we can simply use “update on takes”

```

create function credits_earned() -- the trigger function
returns trigger as $$
begin
    if (new.grade <> 'F' and new.grade is not null
        and (old.grade = 'F' or old.grade is null)) then
        update student
        set tot_cred = tot_cred +
            (select credits from course
             where course.course_id= new.course_id)
        where student.id = new.id;
    end if;
end;$$
language plpgsql;

```

testing ...

```

cr uni=# select * from student where id = '98988';
cr  id      | name   | dept_name | tot_cred
-----+-----+-----+-----
cr  98988   | Tanaka | Biology   |      120
(1 row)

uni=# update takes
uni=# set grade='C+' where id = '98988' and course_id='BIO-301';
UPDATE 1
uni=# select * from student where id = '98988';
cr  id      | name   | dept_name | tot_cred
-----+-----+-----+-----
cr  98988   | Tanaka | Biology   |      124
(1 row)

```

Trigger example

- ❑ **An update** trigger ensuring that amount on account always satisfies $0 \leq \text{amount} \leq 100$

```
create function upd_check() before update on account
returns trigger as $$
begin
    if new.amount < 0 then
        set new.amount = 0;
    elseif new.amount > 100 then
        set new.amount = 100;
    end if;
    return new;
end;

create trigger upd_check_trig -- the trigger
after update on takes for each row execute procedure
upd_check();
```

Statement Level Triggers

- ❑ Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction.
 - Can be more efficient when dealing with SQL statements that update a large number of rows
- ❑ **Supported by** some DBMS' (including **Postgres**) using
 - **for each statement** instead of **for each row**
- ❑ **Insertion of 887000 rows:**

```
insert into movie.movie
select id, title, production_year
from imdb_movie.movie
where kind_id=1;
```
- ❑ **with row-level: 887000 actions, with statement level: 1 action**