IT security

Monday 11th March Course day #5

Theme B: Software and system security

Case: The Code Red buffer overflow attack

Niels Christian Juul (ncjuul@ruc.dk) Niels Jørgensen (nielsj@ruc.dk)

Literature

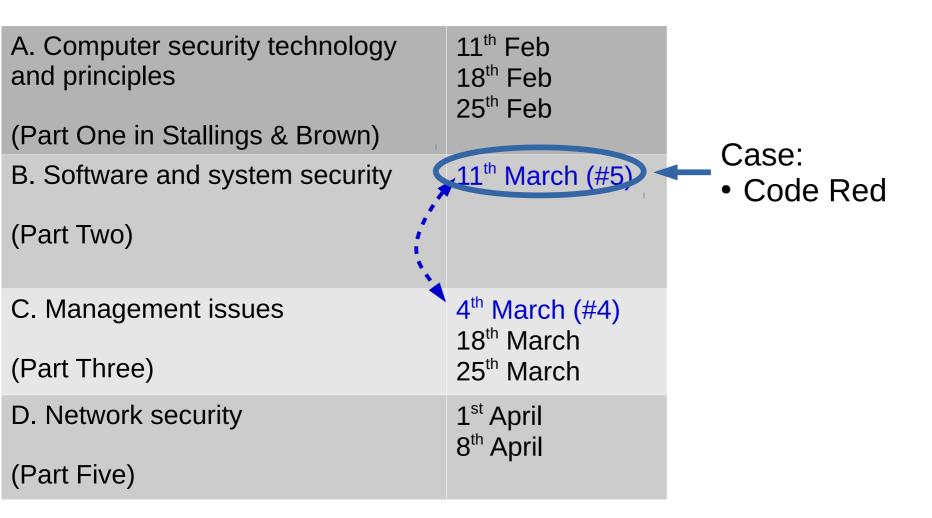
Stallings & Brown

- Chapter 10:
 - buffer overflow: 341-347, 364-367
- Chapter 13:
 - cloud computing: 445-451, 457-464
 - IoT: 466-473

Additional mandatory literature:

- "Code red", Communications of the ACM, December 2001.
- "Webcam hack shows vulnerability of connected devices", Engineering & Technology December 2016.

Four themes (A-D)



Exam questions for theme B (see list on moodle)

Q7: What is a buffer overflow?

Q8: What are the main preventive measures an enterprise can take against buffer overflows and against threats to cloud and IoT security?

Plan for today



Buffer overflows: introduction

Code Red (Anton)

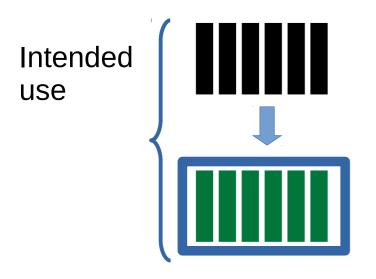
Buffer overflows: how they work

Buffer overflows: prevention

Cloud security

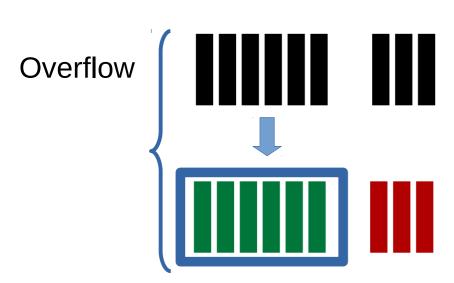
IoT security

Definition of a buffer overflow



".. more input can be placed into a buffer .. than the capacity allocated, overwriting other information"

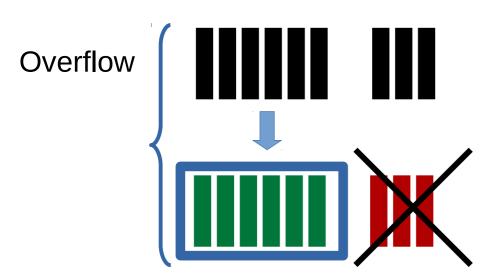
(Stallings & Brown, p 343)



Prevention

A program must reject attempts to write more data to a buffer (an array) than the capacity allocated.

Rejection may be by array bounds checking, as in Java



Exploiting buffer overflows

Exploiting a buffer overflow provides a propagation method.

An exploit may propagate different paylods

- malicious data...
- malicious programs..

Exploiting buffer overflows to run mailicious programs is perhaps the most common type of exploit

- Morris worm
- Code Red

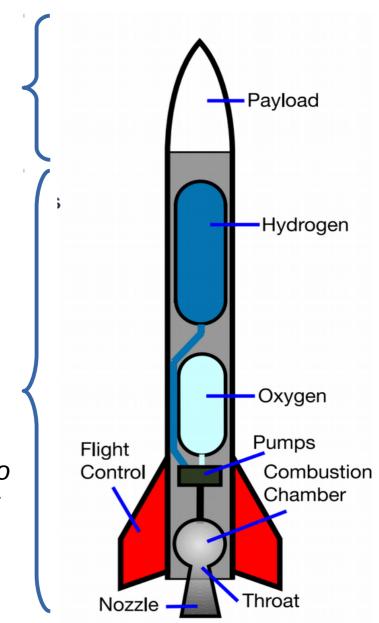
• ..

A malware's payload

• it's impact on the computer

A malware's propagation method

 how it spreads to other computers



Plan for today

Buffer overflows: introduction



Code Red (Anton)

Buffer overflows: how they work

Buffer overflows: prevention

Cloud security

IoT security

Code Red buffer overflow



%u9090%u6858%ucbd3%u7801%u9090%u6858%ucbd3%u7801 %u9090%u6858%ucbd3%u7801%u9090%u9090%u8190%u00c3 %u0003%u8b00%u531b%u53ff%u0078%u0000%u00=a HTTP/1.0

Plan for today

Buffer overflows: introduction

Code Red (Anton)



Buffer overflows: how they work

Buffer overflows: prevention

Cloud security

IoT security

Buffer overflows in C Stallings & Brown, Figure 10.1, p 344

```
int main(int argc, char *argv[]) {
   int valid = FALSE;
   char str1[8];
   char str2[8];

   next_tag(str1);
   gets(str2);
   if (strncmp(str1, str2, 8) == 0)
      valid = TRUE;
   printf("buffer1: str1(%s), str2(%s), valid(%d)\n", str1, str2, valid);
}
```

Suppose we are debugging the program

- therefore, we have inserted the println statement
- Step 1: What is the C program supposed to do?
- Step 2: In what ways is the C program not doing what it is supposed to?
- Step 3: Understand the C program
- Step 4: Why is the C program not doing what it is supposed to?

Step 1: What is the C program supposed to do?

```
int main(int argc, char *argv[]) {
   int valid = FALSE;
   char str1[8];
   char str2[8];

   next_tag(str1);
   gets(str2);
   if (strncmp(str1, str2, 8) == 0)
      valid = TRUE;
   printf("buffer1: str1(%s), str2(%s), valid(%d)\n", str1, str2, valid);
}
```

The program compares two strings (arrays of characters)

- str1 is assigned by the program (eg., "START")
- str2 is supplied by the user

If the strings are equal, the variable *valid* is assigned the value TRUE

- otherwise it is assigned the value FALSE
- note that TRUE is 1 and FALSE is 0

The program is similar to a program that compares two password hashes:

- one hash retrived from the password file
- another hash resulting from hashing a user-supplied password

My variant of buffer1 (on moodle)

```
#include <stdio.h>
                        // defines printf
                        // defines strncmp
#include <string.h>
void next_tag(char[]);
                        // template ("declaration") of function next_tag()
int main(int argc, char * argv[]) {
 int valid = 0;
 char str1[8];
 char str2[8];
 next tag(str1);
 gets(str2);
 if (strncmp(str1, str2, 8) == 0) valid = 1;
 printf("buffer1: str1(%s), str2(%s), valid(%d)\n", str1, str2, valid);
void next tag(char c∏) {
                                               My variant defines next_tag()
 c[0] = 'S';
 c[1] = 'T';
                                                The function definition
 c[2] = 'A';
                                               uses the formal parameter c
 c[3] = 'R';
 c[4] = 'T';
                                               to assign values to the
 c[5] = '\0';
                                               actual parameter str1
```

Step 1: What is the C program supposed to do?

test case	str1 (input / final)	str2 (input / final)	valid (final)
SB#1	START / START	START / START	1
(NJ)	START / START	STOP / STOP	0

In these two test cases, the program does what it is supposed to.

Step 2: In what ways is the C program not doing what it is supposed to?

test case	str1 (input / final)	str2 (input / final)	valid (final)
SB#2	START / TVALUE	EVILINPUTVALUE / EVILINPUTVALUE	0
SB#3	START / BADINPUT	BADINPUTBADINPUT / BADINPUTBADINPUT	1

In test case SB#2:

part of str2 overflows into str1

In test case SB#3

- part of str2 overflows into srt1
- the string comparison succeeds, because the first eight characters are equal

Step 3: Understand the C program

```
int main(int argc, char *argv[]) {
   int valid = FALSE;
   char str1[8];
   char str2[8];

   next_tag(str1);
   gets(str2);
   if (strncmp(str1, str2, 8) == 0)
      valid = TRUE;
   printf("buffer1: str1(%s), str2(%s), valid(%d)\n", str1, str2, valid);
}
```

char str1[8]: declares an array of characters with length 8

char str2[8]: same

next_tag(str1): calls a user-defined function, and passes a pointer to the array str1

gets(str1): calls the system-defined function gets() for reading a string input by the user

strncmp(.., .., 8) compares the first eight characters of the two strings, returns 0 if equal

Step 3: Understand the C program (cont.)

```
int main(int argc, char *argv[]) {
   int valid = FALSE;
   char str1[8];
   char str2[8];

   next_tag(str1);
   gets(str2);
   if (strncmp(str1, str2, 8) == 0)
      valid = TRUE;
   printf("buffer1: str1(%s), str2(%s), valid(%d)\n", str1, str2, valid);
}
```

printf(format, arguments...)

prints format, where:

- ordinary characters are merely printed
 - eg., the ordinary characters "buffer1: str1(" are merely printed
- conversions specifications convert and print the arguments
 - eg., the conversion spec. %s converts and prints str1 (in this case, the conversion is trivial)

printf("str1(%.3s)", str1) would print "str1(STA)" (the first three characters of str1)

Exercise

test case	str1 (input / final)	str2 (input / final)	valid (final)
ex1	START / ?	STARTXX / ?	?
ex2	START / ?	STARTXXXX / ?	?
ex3	START / ?	XXXXXXXXX /	?

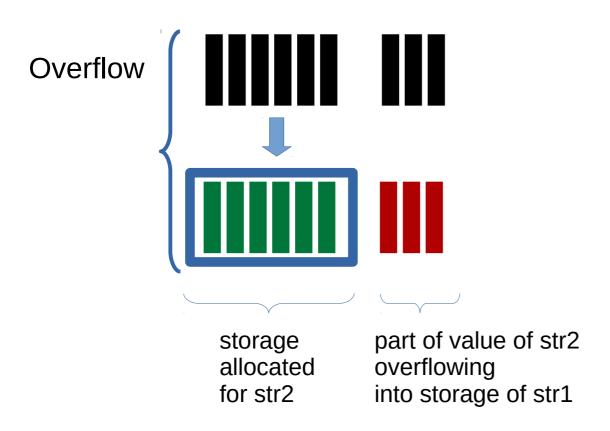
Substitute the question marks with the three final values (as printed)

Exercise solution

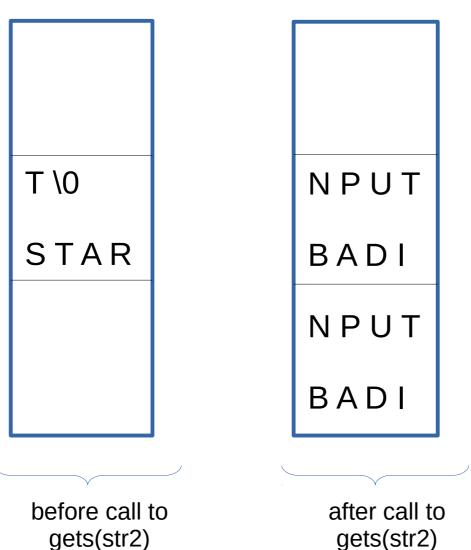
test case	str1 (input / final)	str2 (input / final)	valid (final)
ex1	START / START	STARTXX / STARTXX	0
ex2	START /	STARTXXXX / STARTXXXX	0
ex3	START / XX	XXXXXXXXX / XXXXXXXXXX	0

Step 4: Why is the C program not doing what it is supposed to?

Because the value of str2 overflows into str1



Step 4: Why is the C program not doing what it is supposed to? (cont.)



storage allocated to str1 (8 characters)

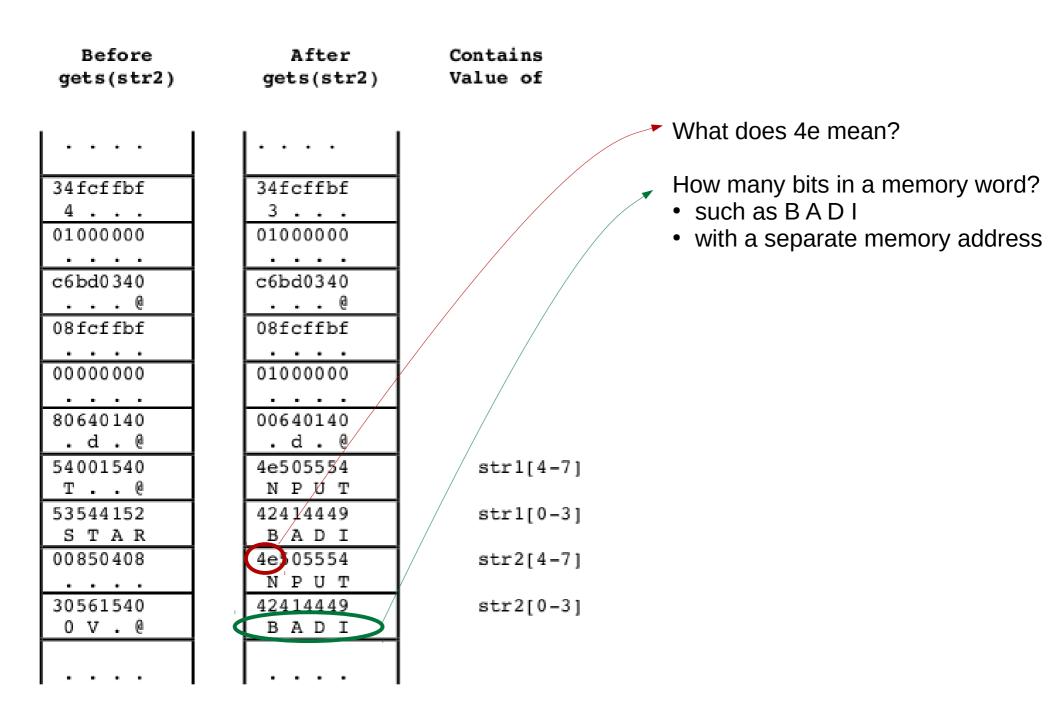
storage allocated to str2 (8 characters)

gets(str2)

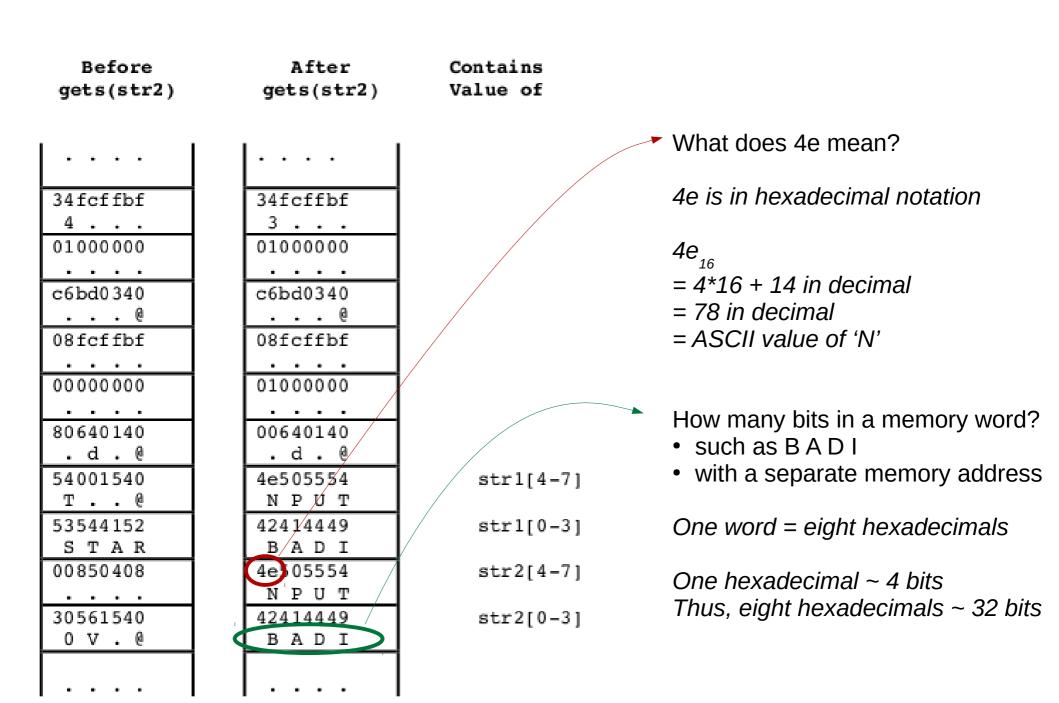
Step 4: details of Fig. 10.2

Before gets(str2)	After gets(str2)	Contains Value of	
34fcffbf 4 01000000 c6bd0340 @ 08fcffbf 00000000 80640140 . d . @	34fcffbf 3 01000000 c6bd0340 @ 08fcffbf 01000000 		 42: hexadecimal number uses 16 symbols 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, a, b, c, d, e, f, 42₁₆ = 4*16 + 2 (in decimal) = 66 (in decimal) = ASCII value of 'B'
54001540 T@	4e505554 N P U T/	str1[4-7]	
53544152 S T A R	42414449 B A D I	str1[0-3]	
00850408	4e505554	str2[4-7]	
30561540 0 V . @	N/PUT 42,14449 BADI	str2[0-3]	

Exercise



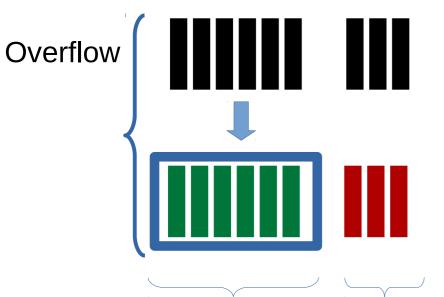
Exercise solution



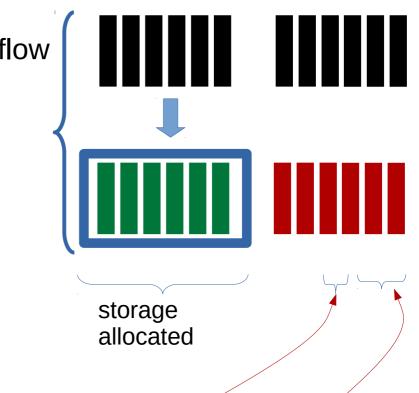
Buffer overflows with malicious code

The program buffer1 contains a buffer overflow which flows into data storage

Buffer overflows may contain malicious code as well



storage allocated for str2 part of value of str2 overflowing into storage of str1



overflow into return address (telling where to continue after end of function)

new return address may point to malicious code

Plan for today

Buffer overflows: introduction

Code Red (Anton)

Buffer overflows: how they work



Buffer overflows: prevention

Cloud security

IoT security

Any internet service must prevent buffer overflows

Clients



Servers

- web
- file servers
- "telephone books"
- •

Exercise

What are the major approaches to prevent buffer overflows?

Exercise solution

What are the major approaches to prevention of buffer overflows?

- 1) Choose a safe programming language
- a language that checks array bounds
- eg., Java
- 2) Hardening of existing programs written in eg., C
- by inspecting and rewriting the programs
- 3) Guidelines for safe use of C
- do not use unsafe system-defined functions, such as gets()
- programmer should always check that sufficient storage is allocated
- •
- 4) Remove library of unsafe system-defined functions

```
import java.util.Scanner;
public class Buffer1 {
 public static void main(String[] args) {
  int valid = 0;
  char[] str1 = new char[8];
  char[] str2 = new char[8];
  next tag(str1);
  gets(str2);
  String s1 = new String(str1);
  String s2 = new String(str2);
  if (s1.equals(s2)) valid = 1;
  System.out.println("buffer1: "
               + "str1(" + s1 + "), "
               + "str2(" + s2 + "), "
               + "valid(" + valid + ")");
 static void next tag(char[] c) {
  c[0] = 'S';
  c[1] = 'T';
  c[2] = 'A';
  c[3] = 'R';
  c[4] = 'T':
 static void gets(char[] c) {
  Scanner scanner = new Scanner(System. in);
  String s = scanner.nextLine();
  for (int i = 0; i < s, length(); i++) {
    c[i] = s.charAt(i):
```

Java has array bounds checking

The execution of

c[i] = s.charAt(i)

aborts if i is 8 or higher, because Java checks i against the array bounds 0 and 7

Exercise

Would it be a good preventive measure against buffer overflows to stop using languages such as C?

Exercise solutions

Would it be a good preventive measure against buffer overflows to stop using languages such as C?

There are also advantages to the use of C

- much more efficient
 - uses less time and space
- more flexible
- most operating systems written in C
 - Unix
 - Mac
 - Windows

Plan for today

Buffer overflows: introduction

Code Red (Anton)

Buffer overflows: how they work

Buffer overflows: prevention



Cloud security

IoT security

Cloud and IoT security





Both cloud and IoT security

- new areas
- with well-known technical solutions
 - encryption (storage, transmission)
 - authentication

Chapter focuses on general concepts

- helps us identify what are main challenges?
- few details, dilemmas, real-life problems

Example: should RUC use cloud services?

RUC's IT systems include:

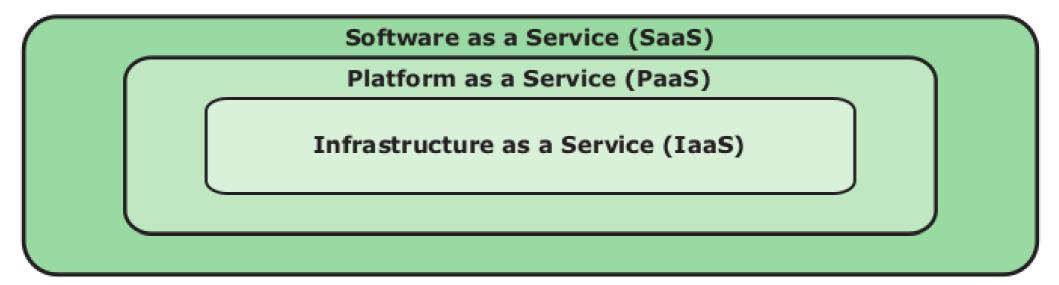
- mail
- moodle (information about courses)
- PCs of employees (including text processing)

RUC's IT department:

- maintain servers and software for mail, moodle, ...
- provide service for employees' PCs

Approximately 40 people

Service models (types of services)



SaaS

client may do word processing, email in cloud

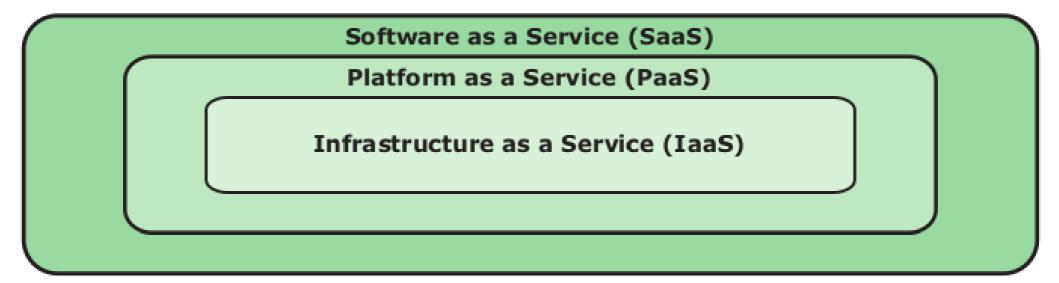
PaaS

client may run specific applications in cloud

IaaS

• client gets access to machines (processing, storage) in cloud

Exercise: what types of services may be relevant to RUC?



SaaS

client may do word processing, email in cloud

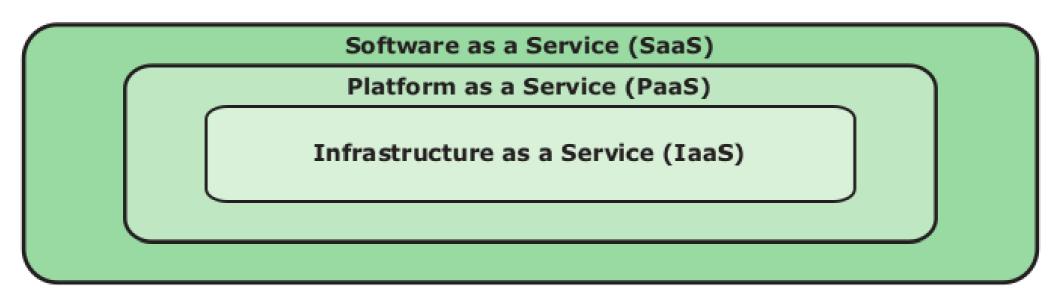
PaaS

client may run specific applications in cloud

IaaS

• client gets access to machines (processing, storage) in cloud

Exercise solution



SaaS

- client may do word processing, email in cloud
- relevant: most RUC employees use word processing and email

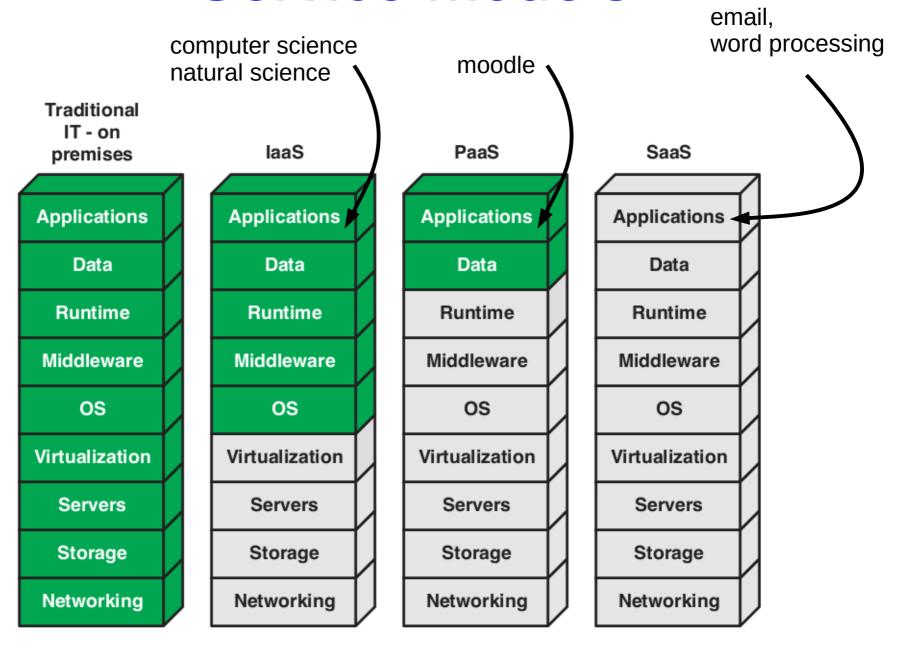
PaaS

- client may run specific applications in cloud
- relevant: RUC runs specific applications such as moodle

laaS

- client gets access to machines (processing, storage) in cloud
- not highly relevant: only employees at computer sc. and natural sc. develop own applications

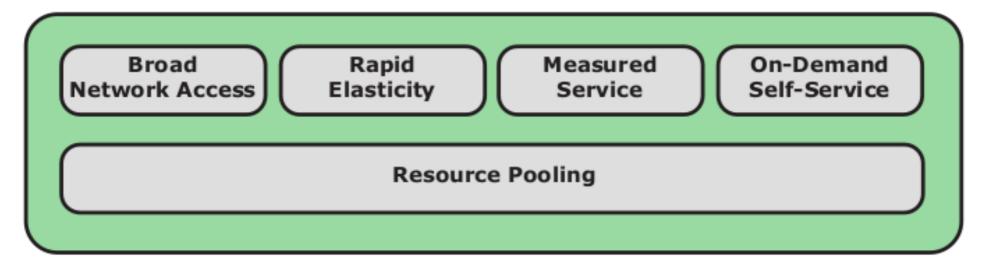
Service models





Managed by cloud service provider

Characteristics (benefits) of cloud computing



Broad network access

RUC employees, students may access resources from laptops, mobiles,...

Rapid elasticity

• moodle scale-up at start of semester, ..

Measured service

(in-house service can be measured/metered just as well)

On-demand self-service, resource pooling

(this is how rapid elasticity is achieved)

Cloud security

".. while the enterprise loses a substantial amount of control over resources, services, and applications, it must maintain accountability for security and privacy policies" (Stallings & Brown, p 457)

CSA categories of cloud security

Stallings & Brown, p 461-462

A category of cloud security is a type of service that may be offered by a cloud service provider (CSP)

Let's run through the categories, with the RUC example in mind.

CSA categories of cloud security

Identity and access management

 CSP must provide access control to word processing and email resourcers

Data loss prevention

- CSP must secure that documents and emails are not lost
- CSC may encrypt documents, emails to prevent leakage

Web security (security of web browsing)

- CSP may provide added protection against, eg. web viruses
- eg., in Adobe Flash Player

E-mail security

- similarly, CSP may provided added protection against email threasts
- eg., spam, phishing

Security assessmes

third party audits

CSA categories of cloud security (cont.)

Exercise: please explain the following categories (possibly by example)

Intrusion management

Security information and event management

Encryption

Business continuity and disaster recovery

Network security

Exercise solution (partial)

Intrusion management

- i. detection, prevention, response
- same as in enterprises own network
- CSP may be more competent/professional

Security information and event management

• monitoring and reporting intrusion and other events

Encryption

• if provided by CSP, entails complex key management because clients must be able to decrypt

Business continuity and disaster recovery

- CSP may provide better backup
- eg., multiple locations

Network security

- firewalls, protection against (D)DoS, ..
- again CSP may be more comptent/professional

Plan for today

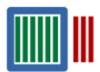
Buffer overflows: introduction

Code Red (Anton)

Buffer overflows: how they work

Buffer overflows: prevention

Cloud security



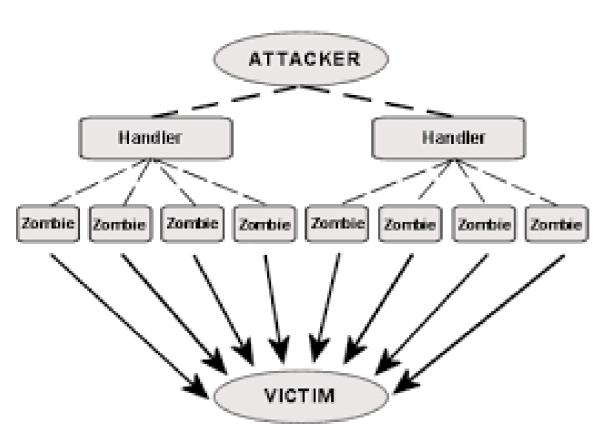
IoT security

IoT security

"IoT is perhaps the most complex and undeveloped area of network security."

Stallings & Brown p 471

The October 2016 webcam hack: a DDoS attack



The hack targetted DYN

 a Internet/DNS service provider (for Twitter and many others)

The hack compromised

- webcams (IoT devices)
- not ordinary computers

Malicious software: Mirai

- attacks linux-based cameras
- worm: scans for more cameras

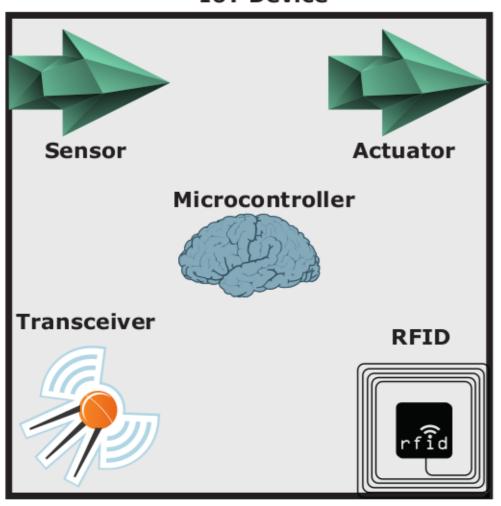
Mirai exploited..

- default usernames/passwords
- some devices can not be updated
- so can not change the default usernames/passwords

IoT devices



IoT Device



Webcam as en IoT example

- sensor
 - the camera
- actuator
 - moving camera angle
- microcontroller
 - may pause camera
- transceiver (internet connection)
 - WIFI
 - IP adress

We will assume webcam has

- unprotected logon
- permits upload/execute programs

How to protect against unprotected IoT devices?

(1) Your own IoT devices

- use IoT devices where you can change the password
- and where software can be updated
- put IoT devices behind a gateway
- this protects your privacy etc.

(2) Your neighbor's IoT devices

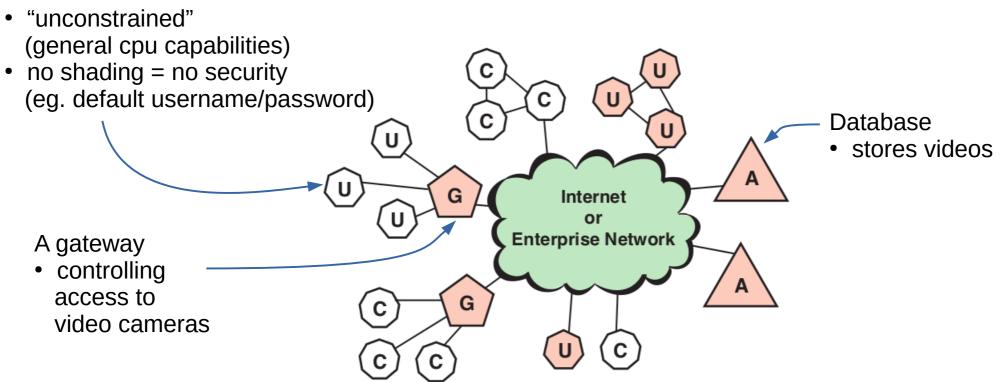
"Even if we purchase better designed products, our neighbours who don't will have their devices hacked and used against the rest of us... The only approach to solving this is .. the designers of these products take their responsibilities seriously".

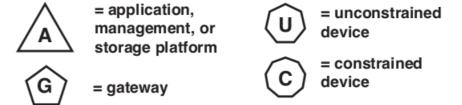
("Webcam hack shows vulnerability of connected devices", p.10).



IoT network

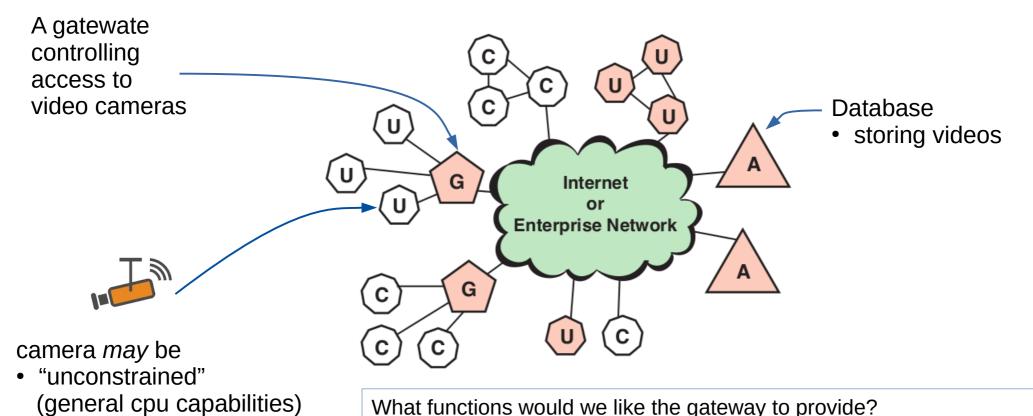
camera assumed to be





shading = includes security features

IoT network



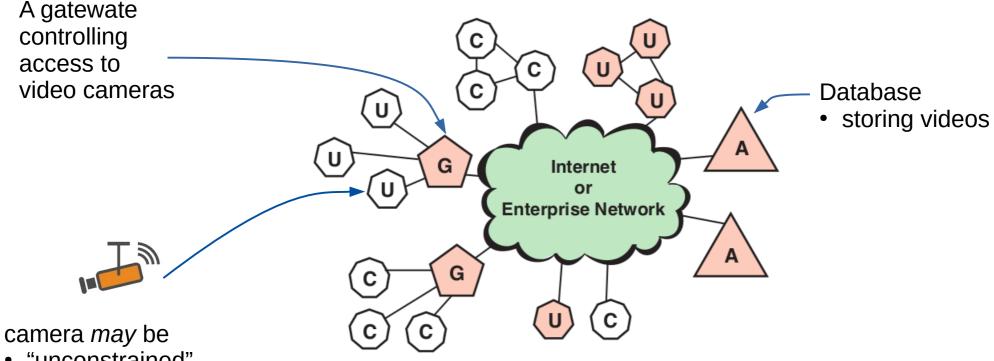
no security

(default username/password)

What functions would we like the gateway to provide? (assuming cameras are unconstrained and have no security)

• see p 472-473 on ITU-T's list of gateway security functions

IoT network



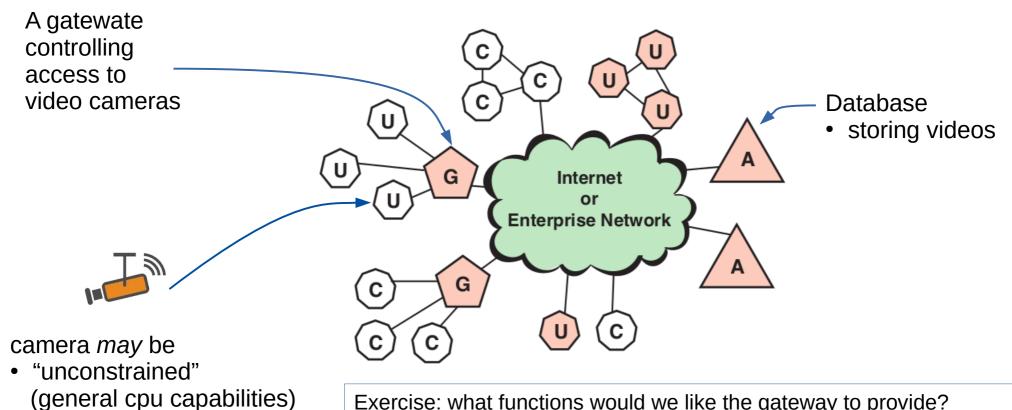
 "unconstrained" (general cpu capabilities)

 no security (default username/password) What functions would we like the gateway to provide? (assuming cameras are unconstrained and have no security)

- identification (logging): who has accessed the cameras?
- access control #1 (authenticate users who access camera)
- access control #2 (authenticate camera) (not highly relevant)

Exercise: please complete the list

Exercise solution



Exercise: what functions would we like the gateway to provide? (assuming camera is unconstrained and without own security)

- identification (logging): who has accessed the cameras?
- access control #1 (authenticate users who access camera)
- access control #2 (authenticate camera) (not highly relevant)
- encrypt data in transmission (we can't encrypt data in the cam.)
- support self-diagnosis (eg., no recording from camera)
- support firm/soft-ware update

no security

(default username/password)

• auto configuration: gateway puts cameras in "pause mode"